

Homework #4

1. Source code

a) Dataloader: I went with the build Dataloader from torch.utils.data for simplicity

```
batch_size = 64

trainloader = torch.utils.data.DataLoader(trainset, batch_size,
                                           shuffle=True, num_workers=2)
valloader = torch.utils.data.DataLoader(validationset, batch_size, num_workers=2)
testloader = torch.utils.data.DataLoader(testset, batch_size,
                                          shuffle=False, num_workers=2)
```

b) Model: Uses additional Conv layers to get better classification and at the end we add a fully connected layer with Dropouts to increase variance in the distribution.

```
device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu') # use GPU
if available

class Net(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv_layer = nn.Sequential(

            # Conv Layer block 1
            nn.Conv2d(in_channels=3, out_channels=32, kernel_size=3, padding=1),
            nn.BatchNorm2d(32),
            nn.ReLU(inplace=True),
            nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2, stride=2),

            # Conv Layer block 2
            nn.Conv2d(in_channels=64, out_channels=128, kernel_size=3, padding=1),
            nn.BatchNorm2d(128),
            nn.ReLU(inplace=True),
            nn.Conv2d(in_channels=128, out_channels=128, kernel_size=3, padding=1),
            nn.ReLU(inplace=True),
            nn.MaxPool2d(kernel_size=2, stride=2),
            nn.Dropout2d(p=0.05),

            # Conv Layer block 3
            nn.Conv2d(in_channels=128, out_channels=256, kernel_size=3, padding=1),
            nn.BatchNorm2d(256),
            nn.ReLU(inplace=True),
            nn.Conv2d(in_channels=256, out_channels=256, kernel_size=3, padding=1),
```

```

        nn.ReLU(inplace=True),
        nn.MaxPool2d(kernel_size=2, stride=2),
    )

    self.fc_layer = nn.Sequential(
        nn.Dropout(p=0.1),
        nn.Linear(4096, 1024),
        nn.ReLU(inplace=True),
        nn.Linear(1024, 512),
        nn.ReLU(inplace=True),
        nn.Dropout(p=0.1),
        nn.Linear(512, 10)
    )

def forward(self, x):
    """Perform forward."""

    # conv layers
    x = self.conv_layer(x)

    # flatten
    x = x.view(x.size(0), -1)

    # fc layer
    x = self.fc_layer(x)

    return x

def training_step(self, batch):
    images, labels = batch[0].to(device), batch[1].to(device)
    out = self(images) # Generate predictions
    loss = F.cross_entropy(out, labels) # Calculate loss
    return loss

@torch.no_grad()
def validation_step(self, batch):
    images, labels = batch[0].to(device), batch[1].to(device)
    out = self(images) # Generate predictions
    loss = F.cross_entropy(out, labels) # Calculate loss
    acc = accuracy(out, labels) # Calculate accuracy
    return {'val_loss': loss.detach(), 'val_acc': acc.detach()}

@torch.no_grad()
def validation_epoch_end(self, outputs):
    batch_losses = [x['val_loss'] for x in outputs]
    epoch_loss = torch.stack(batch_losses).mean() # Combine losses
    batch_accs = [x['val_acc'] for x in outputs]

```

```

        epoch_acc = torch.stack(batch_accs).mean()          # Combine accuracies
        return {'val_loss': epoch_loss.item(), 'val_acc': epoch_acc.item()}

    def epoch_end(self, train_loss, epoch, result):
        print("Epoch [{}], train_loss: {:.4f}, val_loss: {:.4f}, val_acc: {:.4f}".format(epoch, train_loss, result['val_loss'], result['val_acc']))

net = Net()
net = net.to(device)

```

c) Loss Function

```
criterion = nn.CrossEntropyLoss()
```

d) Optimizer

```
optimizer = optim.Adam(net.parameters(), lr=0.001)
```

e) Evaluation

```

# Check the ground truth images
dataiter = iter(testloader)
images, labels = dataiter.next()

imshow(torchvision.utils.make_grid(images[:4]))
print('GroundTruth: ', ' '.join('%5s' % classes[labels[j]] for j in range(4)))

class_correct = list(0. for i in range(10))
class_total = list(0. for i in range(10))
with torch.no_grad():
    for data in testloader:
        images, labels = data[0].to(device), data[1].to(device)
        outputs = net(images)
        _, predicted = torch.max(outputs, 1)
        c = (predicted == labels).squeeze()
        for i in range(4):
            label = labels[i]
            class_correct[label] += c[i].item()
            class_total[label] += 1

for i in range(10):
    print('Accuracy of %5s : %2d %%' % (
        classes[i], 100 * class_correct[i] / class_total[i]))

```

f) Training Loop

```
@torch.no_grad()
def evaluate(model, val_loader):
    outputs = [model.validation_step(batch) for batch in val_loader]
    return model.validation_epoch_end(outputs)

def train(num_epochs, train_loader, val_loader, net, criterion, optimizer):
    history = []
    train_it = 0

    for epoch in range(num_epochs):
        running_loss = 0.0
        for i, data in enumerate(train_loader):
            optimizer.zero_grad()
            # forward
            loss = net.training_step(data)

            # backward
            loss.backward()

            # update the weights
            optimizer.step() # 1 step over optimizer

            running_loss += loss.item()

            train_it += 1

        # Validation phase
        running_loss /= len(train_loader)
        result = evaluate(net, val_loader)
        net.epoch_end(running_loss, epoch, result)
        result['train_loss'] = running_loss
        history.append(result)

    return history
```

g) Datasets and Data augmentation

```
transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

train_transform = transforms.Compose([transforms.RandomCrop(32, padding=4,
padding_mode='reflect'),
                                     transforms.RandomHorizontalFlip(),
                                     transforms.ToTensor(),
                                     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5),
inplace=True)])
```

```
trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                       download=True, transform=train_transform)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                       download=True, transform=transform)

classes = ('plane', 'car', 'bird', 'cat',
          'deer', 'dog', 'frog', 'horse', 'ship', 'truck')


val_size = 10000
train_size = len(trainset) - val_size

trainset, validationset = random_split(trainset, [train_size, val_size])
len(trainset), len(validationset)
```

2 Discussion

2.1 Evolution of training losses and validation losses with standard LeNet using Adam as optimizer, learning rate = 0.001 and 15 epochs

```
Epoch [0], train_loss: 1.7679, val_loss: 1.5919, val_acc: 0.4102
Epoch [1], train_loss: 1.5254, val_loss: 1.4541, val_acc: 0.4551
Epoch [2], train_loss: 1.4113, val_loss: 1.4554, val_acc: 0.4741
Epoch [3], train_loss: 1.3371, val_loss: 1.3202, val_acc: 0.5188
Epoch [4], train_loss: 1.2914, val_loss: 1.2871, val_acc: 0.5285
Epoch [5], train_loss: 1.2551, val_loss: 1.2529, val_acc: 0.5528
Epoch [6], train_loss: 1.2291, val_loss: 1.2044, val_acc: 0.5684
Epoch [7], train_loss: 1.2031, val_loss: 1.2021, val_acc: 0.5692
Epoch [8], train_loss: 1.1826, val_loss: 1.1784, val_acc: 0.5811
Epoch [9], train_loss: 1.1623, val_loss: 1.1772, val_acc: 0.5785
Epoch [10], train_loss: 1.1441, val_loss: 1.1426, val_acc: 0.5989
Epoch [11], train_loss: 1.1302, val_loss: 1.1499, val_acc: 0.5896
Epoch [12], train_loss: 1.1183, val_loss: 1.1505, val_acc: 0.5938
Epoch [13], train_loss: 1.1022, val_loss: 1.1149, val_acc: 0.5993
Epoch [14], train_loss: 1.0912, val_loss: 1.1035, val_acc: 0.6050
```





2.2 Effects of parameter choices

2.2.1 SGD Optimizer

As it can be appreciated in the following images a different optimizer such as SGD in augmented data produces worse results since learning is not as effective as with Adam optimizer




```
Epoch [0], train_loss: 1.3938, val_loss: 1.3751, val_acc: 0.5109
Epoch [1], train_loss: 1.3822, val_loss: 1.3796, val_acc: 0.5045
Epoch [2], train_loss: 1.3795, val_loss: 1.3624, val_acc: 0.5107
Epoch [3], train_loss: 1.3744, val_loss: 1.3667, val_acc: 0.5087
Epoch [4], train_loss: 1.3696, val_loss: 1.3683, val_acc: 0.5101
Epoch [5], train_loss: 1.3679, val_loss: 1.3643, val_acc: 0.5057
Epoch [6], train_loss: 1.3669, val_loss: 1.3616, val_acc: 0.5184
Epoch [7], train_loss: 1.3618, val_loss: 1.3483, val_acc: 0.5206
Epoch [8], train_loss: 1.3580, val_loss: 1.3561, val_acc: 0.5166
Epoch [9], train_loss: 1.3584, val_loss: 1.3498, val_acc: 0.5089
Epoch [10], train_loss: 1.3572, val_loss: 1.3625, val_acc: 0.5097
Epoch [11], train_loss: 1.3508, val_loss: 1.3528, val_acc: 0.5138
Epoch [12], train_loss: 1.3519, val_loss: 1.3527, val_acc: 0.5198
Epoch [13], train_loss: 1.3467, val_loss: 1.3442, val_acc: 0.5190
Epoch [14], train_loss: 1.3509, val_loss: 1.3326, val_acc: 0.5225
```



Overall Accuracy: 55.46%

2.2.2 Using 4x4 filter size instead of 5x5 in Convolutional layers


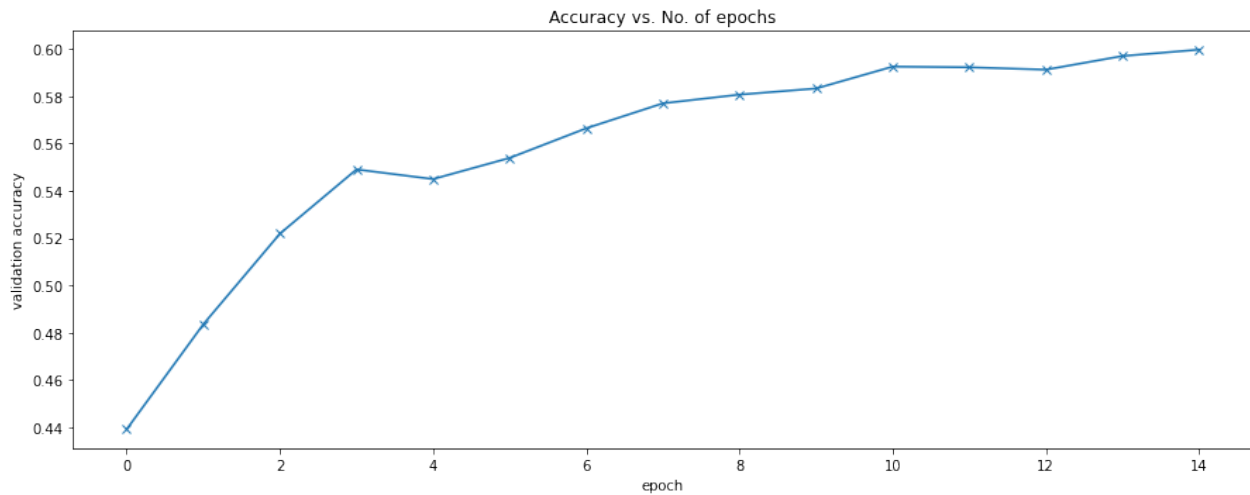
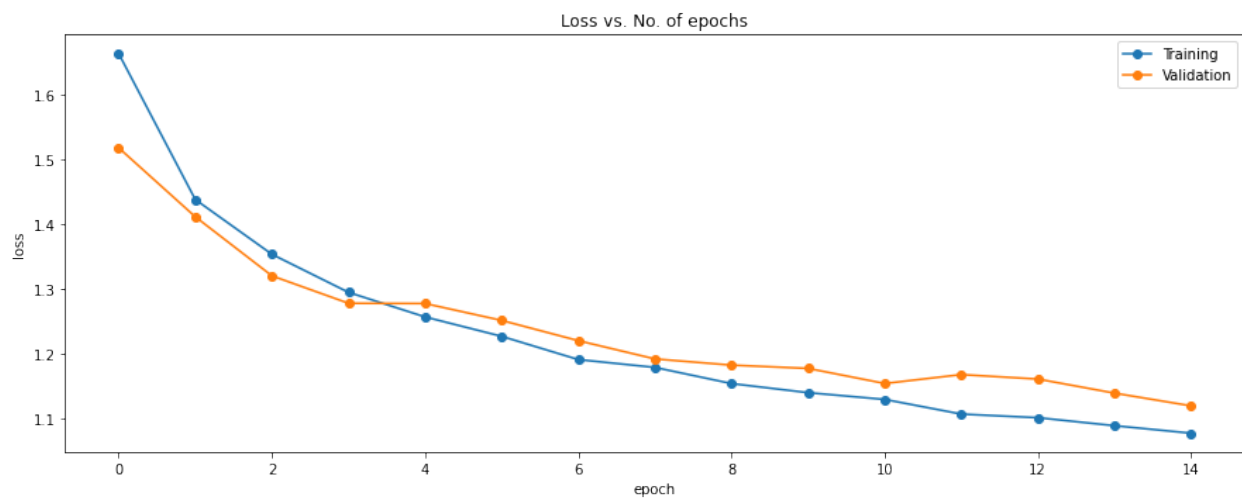
```
Epoch [0], train_loss: 1.8183, val_loss: 1.6362, val_acc: 0.3952
Epoch [1], train_loss: 1.5674, val_loss: 1.4934, val_acc: 0.4547
Epoch [2], train_loss: 1.4569, val_loss: 1.4402, val_acc: 0.4790
Epoch [3], train_loss: 1.3959, val_loss: 1.3747, val_acc: 0.4983
Epoch [4], train_loss: 1.3431, val_loss: 1.3266, val_acc: 0.5211
Epoch [5], train_loss: 1.2973, val_loss: 1.3119, val_acc: 0.5256
Epoch [6], train_loss: 1.2687, val_loss: 1.2681, val_acc: 0.5452
Epoch [7], train_loss: 1.2437, val_loss: 1.2754, val_acc: 0.5404
Epoch [8], train_loss: 1.2258, val_loss: 1.2327, val_acc: 0.5588
Epoch [9], train_loss: 1.2077, val_loss: 1.2171, val_acc: 0.5634
Epoch [10], train_loss: 1.1879, val_loss: 1.2080, val_acc: 0.5634
Epoch [11], train_loss: 1.1703, val_loss: 1.1932, val_acc: 0.5717
Epoch [12], train_loss: 1.1531, val_loss: 1.1466, val_acc: 0.5898
Epoch [13], train_loss: 1.1362, val_loss: 1.1870, val_acc: 0.5740
Epoch [14], train_loss: 1.1191, val_loss: 1.1597, val_acc: 0.5874
```



Accuracy: 63.64%

2.2.3 Using Batch normalization and Leaky_ReLU activation functions



```
Epoch [0], train_loss: 1.6639, val_loss: 1.5182, val_acc: 0.4392
Epoch [1], train_loss: 1.4383, val_loss: 1.4117, val_acc: 0.4836
Epoch [2], train_loss: 1.3537, val_loss: 1.3204, val_acc: 0.5219
Epoch [3], train_loss: 1.2949, val_loss: 1.2780, val_acc: 0.5491
Epoch [4], train_loss: 1.2569, val_loss: 1.2776, val_acc: 0.5450
Epoch [5], train_loss: 1.2269, val_loss: 1.2515, val_acc: 0.5539
Epoch [6], train_loss: 1.1909, val_loss: 1.2202, val_acc: 0.5665
Epoch [7], train_loss: 1.1788, val_loss: 1.1920, val_acc: 0.5770
Epoch [8], train_loss: 1.1538, val_loss: 1.1824, val_acc: 0.5807
Epoch [9], train_loss: 1.1398, val_loss: 1.1773, val_acc: 0.5833
Epoch [10], train_loss: 1.1294, val_loss: 1.1540, val_acc: 0.5926
Epoch [11], train_loss: 1.1067, val_loss: 1.1678, val_acc: 0.5923
Epoch [12], train_loss: 1.1011, val_loss: 1.1608, val_acc: 0.5913
Epoch [13], train_loss: 1.0888, val_loss: 1.1391, val_acc: 0.5970
Epoch [14], train_loss: 1.0772, val_loss: 1.1194, val_acc: 0.5997
```

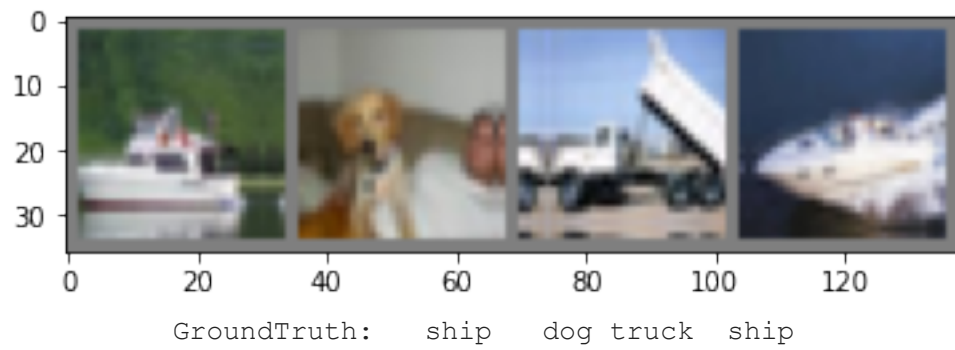


2.2.4 Optimal parameters

In comparison Adam optimizer and learning rate = 0.001 produces much better results than SGD. Moreover, adding convolutional layers, dropout and batch normalization improves accuracy significantly as it can be appreciated in the following results:

```
Epoch [0], train_loss: 1.6764, val_loss: 1.3976, val_acc: 0.4755
Epoch [1], train_loss: 1.2933, val_loss: 1.1780, val_acc: 0.5775
Epoch [2], train_loss: 1.1287, val_loss: 1.0484, val_acc: 0.6246
Epoch [3], train_loss: 1.0226, val_loss: 1.0036, val_acc: 0.6450
Epoch [4], train_loss: 0.9406, val_loss: 0.8996, val_acc: 0.6835
Epoch [5], train_loss: 0.8790, val_loss: 0.8807, val_acc: 0.6936
Epoch [6], train_loss: 0.8361, val_loss: 0.8140, val_acc: 0.7134
Epoch [7], train_loss: 0.7906, val_loss: 0.8205, val_acc: 0.7237
Epoch [8], train_loss: 0.7587, val_loss: 0.7507, val_acc: 0.7377
Epoch [9], train_loss: 0.7215, val_loss: 0.7441, val_acc: 0.7389
Epoch [10], train_loss: 0.6978, val_loss: 0.7449, val_acc: 0.7451
Epoch [11], train_loss: 0.6702, val_loss: 0.7387, val_acc: 0.7478
Epoch [12], train_loss: 0.6487, val_loss: 0.6753, val_acc: 0.7714
Epoch [13], train_loss: 0.6320, val_loss: 0.6648, val_acc: 0.7722
Epoch [14], train_loss: 0.6117, val_loss: 0.6626, val_acc: 0.7763
```






Overall Accuracy: 78.54%

Top-1 prediction accuracy:

Accuracy of plane : 83 %
 Accuracy of car : 88 %
 Accuracy of bird : 65 %
 Accuracy of cat : 58 %
 Accuracy of deer : 65 %
 Accuracy of dog : 67 %
 Accuracy of frog : 85 %
 Accuracy of horse : 71 %
 Accuracy of ship : 94 %
 Accuracy of truck : 89 %

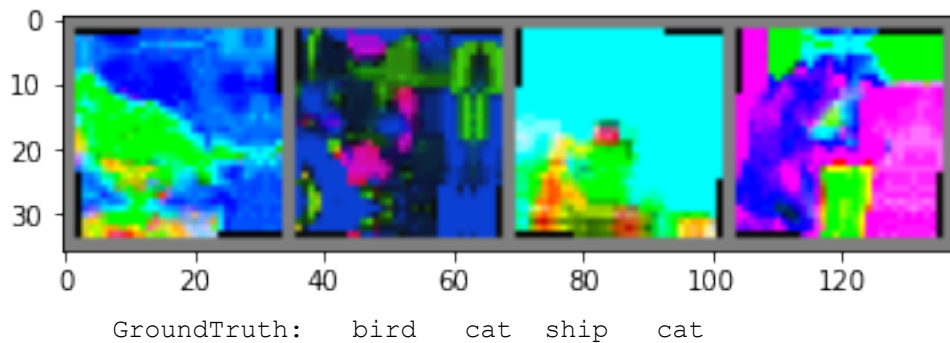
Confusion matrix



[746	19	53	12	13	0	5	8	109	35]	(0)	plane
[9	871	2	1	1	1	5	1	32	77]	(1)	car
[46	2	677	43	52	57	76	25	17	5]	(2)	bird
[10	6	56	615	44	135	67	28	15	24]	(3)	cat
[10	0	56	55	733	14	61	61	5	5]	(4)	deer
[7	2	41	132	34	726	12	36	5	5]	(5)	dog
[6	1	52	41	7	10	867	1	11	4]	(6)	frog
[14	2	27	30	59	60	5	784	7	12]	(7)	horse
[16	16	12	5	4	3	7	4	922	11]	(8)	ship
[14	25	1	3	1	2	3	2	36	913]	(9)	truck

2.3 Example of failed cases

Changing drastically the training data with transformations such as a different hue space or saturation values or random rotations gives above 50% accuracy which shows the strength of Convolutional Neural Networks to work well any type data initialization and regularization.




More important are hyperparameters such as learning rate which can affect learning significantly if the wrong value is chosen.

2.3.1 High learning rate = Fail

As seen in class a low learning rate will take a long time to learn where as a high learning rate can make the gradient either to explode or to jump around without failing in local minima which is what occurs in the following example:

```
Epoch [0], train_loss: 211266.4182, val_loss: 2.3638, val_acc: 0.0999
Epoch [1], train_loss: 2.3686, val_loss: 2.4382, val_acc: 0.0989
Epoch [2], train_loss: 2.3774, val_loss: 2.3666, val_acc: 0.0989
Epoch [3], train_loss: 2.3712, val_loss: 2.3770, val_acc: 0.1028
Epoch [4], train_loss: 2.3667, val_loss: 2.4131, val_acc: 0.1042
Epoch [5], train_loss: 2.3759, val_loss: 2.3973, val_acc: 0.0989
Epoch [6], train_loss: 2.3689, val_loss: 2.3469, val_acc: 0.0989
Epoch [7], train_loss: 2.3724, val_loss: 2.3493, val_acc: 0.1003
Epoch [8], train_loss: 118.2178, val_loss: 2.3806, val_acc: 0.0989
Epoch [9], train_loss: 2.3677, val_loss: 2.4584, val_acc: 0.1028
Epoch [10], train_loss: 2.3771, val_loss: 2.3687, val_acc: 0.0999
Epoch [11], train_loss: 2.3705, val_loss: 2.4444, val_acc: 0.0884
Epoch [12], train_loss: 2.3813, val_loss: 2.5112, val_acc: 0.1028
Epoch [13], train_loss: 2.3778, val_loss: 2.4851, val_acc: 0.1032
Epoch [14], train_loss: 2.3776, val_loss: 2.3423, val_acc: 0.0989
```



This model makes all test cases to fall under one category. It is straightforward to see that it would reach only around 10% accuracy which demonstrates why is important to assign an appropriate learning rate that does not underfit (high learning rate) nor overfit (low learning rate) our model.

